

Apply 3 different CNN's on the MNIST dataset

In [8]:

```
# Credits: https://github.com/keras-team/keras/blob/master/examples/mnist_cnn.py
#Refer this link for making better CNN networks
#https://towardsdatascience.com/a-guide-to-an-efficient-way-to-build-neural-network-architecturespart-ii-hyper-parameter-42efca01e5d7
import warnings
warnings.filterwarnings("ignore")
#from __future__ import print_function
exec('from __future__ import absolute_import, division, print_function')
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K
batch_size = 128
num_classes = 10
epochs = 12
# Preparing training and testing data
# input image dimensions
img_rows, img_cols = 28, 28
# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()
#print(x_train.shape)
if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')
# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
```

Using TensorFlow backend.

```
x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples
```

In [10]:

```
%matplotlib notebook
%matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334 # this function is used to update the plots for
# each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

Model 1-> 2 conv + 2 maxpoll+ 3 dense layers

In [12]:

```

import warnings
warnings.filterwarnings("ignore")
# In this (First Model) lets follow the general structure of the lenet we will make a simple model
# Network Architecture
# input -> conv -> pooling -> conv -> pooling -> FC -> FC -> output
# 8 16 120 84 10
model = Sequential()
model.add(Conv2D(8, kernel_size=(3, 3), activation='relu', padding='same', input_shape=input_shape))
model.add(MaxPooling2D(pool_size=(2, 2), strides=2))# for the location invariants
model.add(Conv2D(16, (5, 5), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2), strides=2))# for the location invariants
model.add(Flatten())
model.add(Dense(120, activation='relu'))
model.add(Dense(84, activation='relu'))
model.add(Dense(num_classes, activation='softmax'))
model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.adam(),
              metrics=['accuracy'])
# this will train the model and validate the model in this fit function
model.summary()

```

Layer (type)	Output Shape	Param #
=====		
conv2d_3 (Conv2D)	(None, 28, 28, 8)	80
max_pooling2d_3 (MaxPooling2D)	(None, 14, 14, 8)	0
conv2d_4 (Conv2D)	(None, 10, 10, 16)	3216
max_pooling2d_4 (MaxPooling2D)	(None, 5, 5, 16)	0
flatten_2 (Flatten)	(None, 400)	0
dense_4 (Dense)	(None, 120)	48120
dense_5 (Dense)	(None, 84)	10164
dense_6 (Dense)	(None, 10)	850
=====		
Total params: 62,430		
Trainable params: 62,430		
Non-trainable params: 0		

In [14]:

```
import warnings
warnings.filterwarnings("ignore")
history=model.fit(x_train, y_train,
                  batch_size=batch_size,
                  epochs=epochs,
                  verbose=1,
                  validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 17s 288us/step - loss: 0.06

89 - acc: 0.9789 - val_loss: 0.0538 - val_acc: 0.9838

Epoch 2/12

60000/60000 [=====] - 15s 243us/step - loss: 0.05

10 - acc: 0.9845 - val_loss: 0.0496 - val_acc: 0.9834

Epoch 3/12

60000/60000 [=====] - 15s 246us/step - loss: 0.03

95 - acc: 0.9879 - val_loss: 0.0364 - val_acc: 0.9880

Epoch 4/12

60000/60000 [=====] - 15s 254us/step - loss: 0.03

18 - acc: 0.9898 - val_loss: 0.0342 - val_acc: 0.9877

Epoch 5/12

60000/60000 [=====] - 15s 258us/step - loss: 0.02

90 - acc: 0.9908 - val_loss: 0.0356 - val_acc: 0.9884

Epoch 6/12

60000/60000 [=====] - 15s 249us/step - loss: 0.02

39 - acc: 0.9921 - val_loss: 0.0292 - val_acc: 0.9912

Epoch 7/12

60000/60000 [=====] - 15s 255us/step - loss: 0.02

15 - acc: 0.9932 - val_loss: 0.0334 - val_acc: 0.9884

Epoch 8/12

60000/60000 [=====] - 15s 252us/step - loss: 0.01

81 - acc: 0.9940 - val_loss: 0.0331 - val_acc: 0.9895

Epoch 9/12

60000/60000 [=====] - 15s 247us/step - loss: 0.01

56 - acc: 0.9948 - val_loss: 0.0322 - val_acc: 0.9898

Epoch 10/12

60000/60000 [=====] - 15s 249us/step - loss: 0.01

33 - acc: 0.9956 - val_loss: 0.0348 - val_acc: 0.9894

Epoch 11/12

60000/60000 [=====] - 15s 252us/step - loss: 0.01

25 - acc: 0.9957 - val_loss: 0.0407 - val_acc: 0.9869

Epoch 12/12

60000/60000 [=====] - 15s 251us/step - loss: 0.01

11 - acc: 0.9965 - val_loss: 0.0391 - val_acc: 0.9881

Test loss: 0.039050305695623684

Test accuracy: 0.9881

In [15]:


```
score = model.evaluate(x_train, y_train, verbose=0)
print('Train score:', score[0])
print('Train accuracy:', score[1]*100)
print('\n***** \n')
#test accuracy
score = model.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1]*100)
# plot
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch');
ax.set_ylabel('Categorical Crossentropy Loss')
x = list(range(1,12+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Train score: 0.010427321565627062

Train accuracy: 99.67333333333333

Test score: 0.039050305695623684

Test accuracy: 98.81



Model 2-> 3 conv + 3 maxpoll+ 2 dense layers

In [16]:

```
import warnings
warnings.filterwarnings("ignore")
# go basic model to deep layer model
# Network Architecture
# input -> conv -> polling -> conv -> polling -> conv -> polling -> FC -> output
# 8 32 128 64
model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=input_shape))
model.add(MaxPooling2D(pool_size=(2, 2), strides=2))# for the location invariants
model.add(Conv2D(64, (3,3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2), strides=2))# for the location invariants
model.add(Conv2D(128, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2), strides=2))# for the location invariants
model.add(Flatten())
model.add(Dense(64, activation='relu'))
model.add(Dense(num_classes, activation='softmax'))
model.compile(loss=keras.losses.categorical_crossentropy,
optimizer=keras.optimizers.adam(),
metrics=['accuracy'])
# this will train the model and validate the model in this fit function
model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_5 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_5 (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_6 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_6 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_7 (Conv2D)	(None, 3, 3, 128)	73856
max_pooling2d_7 (MaxPooling2D)	(None, 1, 1, 128)	0
flatten_3 (Flatten)	(None, 128)	0
dense_7 (Dense)	(None, 64)	8256
dense_8 (Dense)	(None, 10)	650
Total params: 101,578		
Trainable params: 101,578		
Non-trainable params: 0		

In [17]:

```
import warnings
warnings.filterwarnings("ignore")
history=model.fit(x_train, y_train,
                  batch_size=batch_size,
                  epochs=epochs,
                  verbose=1,
                  validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 40s 670us/step - loss: 0.32

16 - acc: 0.9046 - val_loss: 0.1012 - val_acc: 0.9684

Epoch 2/12

60000/60000 [=====] - 38s 627us/step - loss: 0.09

35 - acc: 0.9714 - val_loss: 0.0711 - val_acc: 0.9772

Epoch 3/12

60000/60000 [=====] - 37s 620us/step - loss: 0.06

65 - acc: 0.9797 - val_loss: 0.0708 - val_acc: 0.9789

Epoch 4/12

60000/60000 [=====] - 37s 620us/step - loss: 0.05

37 - acc: 0.9834 - val_loss: 0.0510 - val_acc: 0.9845

Epoch 5/12

60000/60000 [=====] - 38s 627us/step - loss: 0.04

58 - acc: 0.9854 - val_loss: 0.0579 - val_acc: 0.9832

Epoch 6/12

60000/60000 [=====] - 37s 620us/step - loss: 0.03

80 - acc: 0.9881 - val_loss: 0.0448 - val_acc: 0.9861

Epoch 7/12

60000/60000 [=====] - 37s 624us/step - loss: 0.03

23 - acc: 0.9901 - val_loss: 0.0474 - val_acc: 0.9864

Epoch 8/12

60000/60000 [=====] - 38s 626us/step - loss: 0.02

71 - acc: 0.9917 - val_loss: 0.0484 - val_acc: 0.9867

Epoch 9/12

60000/60000 [=====] - 38s 634us/step - loss: 0.02

31 - acc: 0.9926 - val_loss: 0.0564 - val_acc: 0.9846

Epoch 10/12

60000/60000 [=====] - 37s 622us/step - loss: 0.02

13 - acc: 0.9931 - val_loss: 0.0584 - val_acc: 0.9823

Epoch 11/12

60000/60000 [=====] - 38s 627us/step - loss: 0.01

74 - acc: 0.9946 - val_loss: 0.0461 - val_acc: 0.9872

Epoch 12/12

60000/60000 [=====] - 38s 629us/step - loss: 0.01

62 - acc: 0.9948 - val_loss: 0.0454 - val_acc: 0.9881

Test loss: 0.04539995787585067

Test accuracy: 0.9881

In [18]:


```
score = model.evaluate(x_train, y_train, verbose=0)
print('Train score:', score[0])
print('Train accuracy:', score[1]*100)
print('\n***** \n')
#test accuracy
score = model.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1]*100)
# plot
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch');
ax.set_ylabel('Categorical Crossentropy Loss')
x = list(range(1,12+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Train score: 0.009743525346997921

Train accuracy: 99.69833333333334

Test score: 0.04539995787585067

Test accuracy: 98.81



Finally we train a model with the trend Conv-Conv-Pool-Conv-Conv-Pool

Model 3 -> 4 conv+ 2 maxpoll + 2 dence

In [19]:

```
# go basic model to deep layer model
# Network Architecture
# input -> conv -> conv -> polling -> conv -> conv -> polling -> FC -> output
# 16 16 32 32 512
model = Sequential()
model.add(Conv2D(16, kernel_size=(3, 3), activation='relu', padding='same', input_shape=input_shape))
model.add(Conv2D(16, (3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D(pool_size=(2, 2), strides=2)) # for the location invariants
model.add(Conv2D(32, (3, 3), activation='relu'))
model.add(Conv2D(32, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2), strides=2)) # for the location invariants
model.add(Flatten())
model.add(Dense(512, activation='relu'))
model.add(Dense(num_classes, activation='softmax'))
model.compile(loss=keras.losses.categorical_crossentropy,
optimizer=keras.optimizers.adam(),
metrics=['accuracy'])
# this will train the model and validate the model in this fit function
model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_8 (Conv2D)	(None, 28, 28, 16)	160
conv2d_9 (Conv2D)	(None, 28, 28, 16)	2320
max_pooling2d_8 (MaxPooling2D)	(None, 14, 14, 16)	0
conv2d_10 (Conv2D)	(None, 12, 12, 32)	4640
conv2d_11 (Conv2D)	(None, 10, 10, 32)	9248
max_pooling2d_9 (MaxPooling2D)	(None, 5, 5, 32)	0
flatten_4 (Flatten)	(None, 800)	0
dense_9 (Dense)	(None, 512)	410112
dense_10 (Dense)	(None, 10)	5130
Total params: 431,610		
Trainable params: 431,610		
Non-trainable params: 0		

In [20]:

```
import warnings
warnings.filterwarnings("ignore")
history=model.fit(x_train, y_train,
                  batch_size=batch_size,
                  epochs=epochs,
                  verbose=1,
                  validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 64s 1ms/step - loss: 0.1727
- acc: 0.9479 - val_loss: 0.0515 - val_acc: 0.9826

Epoch 2/12

60000/60000 [=====] - 59s 987us/step - loss: 0.04
42 - acc: 0.9862 - val_loss: 0.0300 - val_acc: 0.9905

Epoch 3/12

60000/60000 [=====] - 59s 983us/step - loss: 0.03
06 - acc: 0.9906 - val_loss: 0.0258 - val_acc: 0.9912

Epoch 4/12

60000/60000 [=====] - 61s 1ms/step - loss: 0.0221
- acc: 0.9931 - val_loss: 0.0296 - val_acc: 0.9905

Epoch 5/12

60000/60000 [=====] - 67s 1ms/step - loss: 0.0173
- acc: 0.9940 - val_loss: 0.0294 - val_acc: 0.9912

Epoch 6/12

60000/60000 [=====] - 57s 957us/step - loss: 0.01
47 - acc: 0.9952 - val_loss: 0.0319 - val_acc: 0.9904

Epoch 7/12

60000/60000 [=====] - 61s 1ms/step - loss: 0.0112
- acc: 0.9963 - val_loss: 0.0282 - val_acc: 0.9922

Epoch 8/12

60000/60000 [=====] - 58s 963us/step - loss: 0.00
95 - acc: 0.9970 - val_loss: 0.0339 - val_acc: 0.9897

Epoch 9/12

60000/60000 [=====] - 58s 973us/step - loss: 0.00
99 - acc: 0.9967 - val_loss: 0.0374 - val_acc: 0.9897

Epoch 10/12

60000/60000 [=====] - 59s 976us/step - loss: 0.00
70 - acc: 0.9975 - val_loss: 0.0287 - val_acc: 0.9914

Epoch 11/12

60000/60000 [=====] - 59s 980us/step - loss: 0.00
67 - acc: 0.9978 - val_loss: 0.0307 - val_acc: 0.9912

Epoch 12/12

60000/60000 [=====] - 59s 987us/step - loss: 0.00
51 - acc: 0.9986 - val_loss: 0.0242 - val_acc: 0.9935

Test loss: 0.024233603691490725

Test accuracy: 0.9935

In [21]:


```
score = model.evaluate(x_train, y_train, verbose=0)
print('Train score:', score[0])
print('Train accuracy:', score[1]*100)
print('\n***** \n')
#test accuracy
score = model.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1]*100)
# plot
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch');
ax.set_ylabel('Categorical Crossentropy Loss')
x = list(range(1,12+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Train score: 0.0032476350627371935

Train accuracy: 99.90833333333333

Test score: 0.024233603691490725

Test accuracy: 99.35000000000001



Model 1-> 2 conv + 2 maxpoll+ 3 dense layer +Dropout (0.5)

In [23]:

```
#Same models with Dropouts
import warnings
warnings.filterwarnings("ignore")
# In this (First Model) Lets follow the general structure of the lenet we will make a simple model
# Network Architecture
# input -> conv -> pooling -> conv -> pooling -> dropout -> FC -> FC -> output
# 8 16 120 84 10
model = Sequential()
model.add(Conv2D(8, kernel_size=(3, 3), activation='relu', padding='same', input_shape=input_shape))
model.add(MaxPooling2D(pool_size=(2, 2), strides=2))# for the location invariants
model.add(Conv2D(16, (5, 5), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2), strides=2))# for the location invariants
model.add(Dropout(0.5))
model.add(Flatten())
model.add(Dense(120, activation='relu'))
model.add(Dense(84, activation='relu'))
model.add(Dense(num_classes, activation='softmax'))
model.compile(loss=keras.losses.categorical_crossentropy,
optimizer=keras.optimizers.adam(),
metrics=['accuracy'])
# this will train the model and validate the model in this fit function
model.summary()
```

Layer (type)	Output Shape	Param #
=====		
conv2d_14 (Conv2D)	(None, 28, 28, 8)	80
max_pooling2d_12 (MaxPooling)	(None, 14, 14, 8)	0
conv2d_15 (Conv2D)	(None, 10, 10, 16)	3216
max_pooling2d_13 (MaxPooling)	(None, 5, 5, 16)	0
dropout_2 (Dropout)	(None, 5, 5, 16)	0
flatten_6 (Flatten)	(None, 400)	0
dense_14 (Dense)	(None, 120)	48120
dense_15 (Dense)	(None, 84)	10164
dense_16 (Dense)	(None, 10)	850
=====		
Total params: 62,430		
Trainable params: 62,430		
Non-trainable params: 0		

In [24]:

```

history=model.fit(x_train, y_train,
                  batch_size=batch_size,
                  epochs=epochs,
                  verbose=1,
                  validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 17s 284us/step - loss: 0.42
44 - acc: 0.8639 - val_loss: 0.0859 - val_acc: 0.9748

Epoch 2/12

60000/60000 [=====] - 16s 268us/step - loss: 0.13
36 - acc: 0.9576 - val_loss: 0.0558 - val_acc: 0.9824

Epoch 3/12

60000/60000 [=====] - 17s 278us/step - loss: 0.10
68 - acc: 0.9662 - val_loss: 0.0446 - val_acc: 0.9849

Epoch 4/12

60000/60000 [=====] - 17s 276us/step - loss: 0.08
97 - acc: 0.9710 - val_loss: 0.0395 - val_acc: 0.9867

Epoch 5/12

60000/60000 [=====] - 17s 286us/step - loss: 0.07
77 - acc: 0.9746 - val_loss: 0.0351 - val_acc: 0.9891

Epoch 6/12

60000/60000 [=====] - 17s 278us/step - loss: 0.07
33 - acc: 0.9768 - val_loss: 0.0346 - val_acc: 0.9881

Epoch 7/12

60000/60000 [=====] - 17s 277us/step - loss: 0.06
53 - acc: 0.9791 - val_loss: 0.0342 - val_acc: 0.9888

Epoch 8/12

60000/60000 [=====] - 17s 277us/step - loss: 0.06
15 - acc: 0.9804 - val_loss: 0.0302 - val_acc: 0.9898

Epoch 9/12

60000/60000 [=====] - 17s 287us/step - loss: 0.05
72 - acc: 0.9813 - val_loss: 0.0303 - val_acc: 0.9894

Epoch 10/12

60000/60000 [=====] - 17s 282us/step - loss: 0.05
53 - acc: 0.9819 - val_loss: 0.0302 - val_acc: 0.9897

Epoch 11/12

60000/60000 [=====] - 17s 286us/step - loss: 0.05
24 - acc: 0.9833 - val_loss: 0.0260 - val_acc: 0.9910

Epoch 12/12

60000/60000 [=====] - 17s 286us/step - loss: 0.05
00 - acc: 0.9837 - val_loss: 0.0273 - val_acc: 0.9910

Test loss: 0.027327572450373556


Test accuracy: 0.991

In [25]:

```
score = model.evaluate(x_train, y_train, verbose=0)
print('Train score:', score[0])
print('Train accuracy:', score[1]*100)
print('\n***** \n')
#test accuracy
score = model.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1]*100)
# plot
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch');
ax.set_ylabel('Categorical Crossentropy Loss')
x = list(range(1,12+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Train score: 0.018932762710633686
Train accuracy: 99.43666666666667

Test score: 0.027327572450373556
Test accuracy: 99.1



Model 2-> 3 conv + 3 maxpoll+ 2 dense layers + Dropout (0.9)

In [27]:

```
import warnings
warnings.filterwarnings("ignore")
# go basic model to deep layer model
# Network Architecture
# input -> conv -> pooling -> conv -> pooling -> conv -> pooling -> dropout -> FC -> output
# 8 32 128 64
model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=input_shape))
model.add(MaxPooling2D(pool_size=(2, 2), strides=2)) # for the location invariants
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2), strides=2)) # for the location invariants
model.add(Conv2D(128, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2), strides=2)) # for the location invariants
model.add(Dropout(0.9))
model.add(Flatten())
model.add(Dense(64, activation='relu'))
model.add(Dense(num_classes, activation='softmax'))
model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.adam(),
              metrics=['accuracy'])
# this will train the model and validate the model in this fit function
model.summary()
```

W0629 22:43:29.693563 692 nn_ops.py:4224] Large dropout rate: 0.9 (>0.5). In TensorFlow 2.x, dropout() uses dropout rate instead of keep_prob. Please ensure that this is intended.

Layer (type)	Output Shape	Param #
=====		
conv2d_19 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_17 (MaxPooling)	(None, 13, 13, 32)	0
conv2d_20 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_18 (MaxPooling)	(None, 5, 5, 64)	0
conv2d_21 (Conv2D)	(None, 3, 3, 128)	73856
max_pooling2d_19 (MaxPooling)	(None, 1, 1, 128)	0
dropout_4 (Dropout)	(None, 1, 1, 128)	0
flatten_8 (Flatten)	(None, 128)	0
dense_19 (Dense)	(None, 64)	8256
dense_20 (Dense)	(None, 10)	650
=====		
Total params: 101,578		
Trainable params: 101,578		
Non-trainable params: 0		

In [28]:

```
history=model.fit(x_train, y_train,
batch_size=batch_size,
epochs=epochs,
verbose=1,
validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 39s 654us/step - loss: 1.30

13 - acc: 0.5210 - val_loss: 0.2997 - val_acc: 0.9407

Epoch 2/12

60000/60000 [=====] - 37s 615us/step - loss: 0.80

30 - acc: 0.7029 - val_loss: 0.1820 - val_acc: 0.9527

Epoch 3/12

60000/60000 [=====] - 37s 616us/step - loss: 0.67

01 - acc: 0.7518 - val_loss: 0.1444 - val_acc: 0.9614

Epoch 4/12

60000/60000 [=====] - 36s 605us/step - loss: 0.60

75 - acc: 0.7772 - val_loss: 0.1168 - val_acc: 0.9679

Epoch 5/12

60000/60000 [=====] - 37s 620us/step - loss: 0.55

34 - acc: 0.8004 - val_loss: 0.1155 - val_acc: 0.9679

Epoch 6/12

60000/60000 [=====] - 37s 610us/step - loss: 0.51

37 - acc: 0.8165 - val_loss: 0.1043 - val_acc: 0.9701

Epoch 7/12

60000/60000 [=====] - 37s 617us/step - loss: 0.48

18 - acc: 0.8267 - val_loss: 0.0927 - val_acc: 0.9736

Epoch 8/12

60000/60000 [=====] - 37s 613us/step - loss: 0.45

95 - acc: 0.8352 - val_loss: 0.1023 - val_acc: 0.9700

Epoch 9/12

60000/60000 [=====] - 36s 596us/step - loss: 0.43

43 - acc: 0.8478 - val_loss: 0.0919 - val_acc: 0.9736

Epoch 10/12

60000/60000 [=====] - 36s 604us/step - loss: 0.41

74 - acc: 0.8532 - val_loss: 0.0951 - val_acc: 0.9716

Epoch 11/12

60000/60000 [=====] - 36s 599us/step - loss: 0.39

85 - acc: 0.8642 - val_loss: 0.0918 - val_acc: 0.9724

Epoch 12/12

60000/60000 [=====] - 36s 594us/step - loss: 0.38

30 - acc: 0.8711 - val_loss: 0.0968 - val_acc: 0.9701

Test loss: 0.09679163726270198

Test accuracy: 0.9701

In [29]:


```
score = model.evaluate(x_train, y_train, verbose=0)
print('Train score:', score[0])
print('Train accuracy:', score[1]*100)
print('\n***** \n')
#test accuracy
score = model.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1]*100)
# plot
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch');
ax.set_ylabel('Categorical Crossentropy Loss')
x = list(range(1,12+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Train score: 0.08697761877303323

Train accuracy: 97.515

Test score: 0.09679163726270198

Test accuracy: 97.00999999999999



Model 3-> 4 conv + 2 maxpoll+ 2 dense layers + Dropout (0.3)

In [30]:

```
# go basic model to deep layer model
# Network Architecture
# input -> conv -> conv -> polling -> conv -> conv -> polling -> dropout -> FC -> output
# 16 16 32 32 512
model = Sequential()
model.add(Conv2D(16, kernel_size=(3, 3), activation='relu', padding='same', input_shape=input_shape))
model.add(Conv2D(16, (3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D(pool_size=(2, 2), strides=2)) # for the location invariants
model.add(Conv2D(32, (3, 3), activation='relu'))
model.add(Conv2D(32, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2), strides=2)) # for the location invariants
model.add(Dropout(0.3))
model.add(Flatten())
model.add(Dense(512, activation='relu'))
model.add(Dense(num_classes, activation='softmax'))
model.compile(loss=keras.losses.categorical_crossentropy,
optimizer=keras.optimizers.adam(),
metrics=['accuracy'])
# this will train the model and validate the model in this fit function
model.summary()
```

Layer (type)	Output Shape	Param #
=====		
conv2d_22 (Conv2D)	(None, 28, 28, 16)	160
conv2d_23 (Conv2D)	(None, 28, 28, 16)	2320
max_pooling2d_20 (MaxPooling)	(None, 14, 14, 16)	0
conv2d_24 (Conv2D)	(None, 12, 12, 32)	4640
conv2d_25 (Conv2D)	(None, 10, 10, 32)	9248
max_pooling2d_21 (MaxPooling)	(None, 5, 5, 32)	0
dropout_5 (Dropout)	(None, 5, 5, 32)	0
flatten_9 (Flatten)	(None, 800)	0
dense_21 (Dense)	(None, 512)	410112
dense_22 (Dense)	(None, 10)	5130
=====		
Total params: 431,610		
Trainable params: 431,610		
Non-trainable params: 0		

In [31]:

```
history=model.fit(x_train, y_train,
batch_size=batch_size,
epochs=epochs,
verbose=1,
validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 59s 975us/step - loss: 0.20
71 - acc: 0.9347 - val_loss: 0.0499 - val_acc: 0.9845

Epoch 2/12

60000/60000 [=====] - 57s 957us/step - loss: 0.06
00 - acc: 0.9814 - val_loss: 0.0342 - val_acc: 0.9890

Epoch 3/12

60000/60000 [=====] - 58s 963us/step - loss: 0.04
14 - acc: 0.9868 - val_loss: 0.0331 - val_acc: 0.9896

Epoch 4/12

60000/60000 [=====] - 57s 957us/step - loss: 0.03
30 - acc: 0.9895 - val_loss: 0.0222 - val_acc: 0.9932

Epoch 5/12

60000/60000 [=====] - 58s 967us/step - loss: 0.02
86 - acc: 0.9907 - val_loss: 0.0241 - val_acc: 0.9920

Epoch 6/12

60000/60000 [=====] - 59s 979us/step - loss: 0.02
42 - acc: 0.9922 - val_loss: 0.0213 - val_acc: 0.9929

Epoch 7/12

60000/60000 [=====] - 59s 976us/step - loss: 0.01
98 - acc: 0.9939 - val_loss: 0.0239 - val_acc: 0.9923

Epoch 8/12

60000/60000 [=====] - 58s 968us/step - loss: 0.01
82 - acc: 0.9936 - val_loss: 0.0235 - val_acc: 0.9927

Epoch 9/12

60000/60000 [=====] - 58s 965us/step - loss: 0.01
72 - acc: 0.9942 - val_loss: 0.0228 - val_acc: 0.9917

Epoch 10/12

60000/60000 [=====] - 57s 958us/step - loss: 0.01
45 - acc: 0.9953 - val_loss: 0.0279 - val_acc: 0.9915

Epoch 11/12

60000/60000 [=====] - 58s 964us/step - loss: 0.01
39 - acc: 0.9954 - val_loss: 0.0217 - val_acc: 0.9925

Epoch 12/12

60000/60000 [=====] - 58s 959us/step - loss: 0.01
15 - acc: 0.9962 - val_loss: 0.0234 - val_acc: 0.9936

Test loss: 0.023363230100554574

Test accuracy: 0.9936

In [32]:


```
score = model.evaluate(x_train, y_train, verbose=0)
print('Train score:', score[0])
print('Train accuracy:', score[1]*100)
print('\n***** \n')
#test accuracy
score = model.evaluate(x_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1]*100)
# plot
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch');
ax.set_ylabel('Categorical Crossentropy Loss')
x = list(range(1,12+1))
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Train score: 0.00590441487302293

Train accuracy: 99.80166666666666

Test score: 0.023363230100554574

Test accuracy: 99.36



Compare all the model results

In [33]:

```

from prettytable import PrettyTable
tb = PrettyTable()
tb.field_names= ("conv_layers", "MAxPoll_layers", "Dense_layers", "Dropout", "Accuracy")
tb.add_row(["2", "2", "3", "NO", 98.81])
tb.add_row(["3", "3", "2", "NO", 98.81])
tb.add_row(["4", "2", "2", "NO", 99.35])
tb.add_row(["2", "2", "3", "0.5", 99.1])
tb.add_row(["3", "3", "2", "0.9", 97.0])
tb.add_row(["4", "2", "2", "0.3", 99.36])
print(tb.get_string(titles = "CNN Models - Observations"))

```

conv_layers	MAxPoll_layers	Dense_layers	Dropout	Accuracy
2	2	3	NO	98.81
3	3	2	NO	98.81
4	2	2	NO	99.35
2	2	3	0.5	99.1
3	3	2	0.9	97.0
4	2	2	0.3	99.36